

# DDA3020 Machine Learning

## Lecture 05 Linear Regression

Jicong Fan  
School of Data Science, CUHK-SZ

09/18/2025

# Outline

- ① Notations, vectors, matrices
- ② Functions, derivative and gradient
- ③ Modeling of linear regression
  - Deterministic perspective
  - Probabilistic perspective
- ④ Learning of linear regression
  - Analytical solution
  - Gradient descent algorithm
- ⑤ Linear regression of multiple outputs
- ⑥ Linear regression for classification
- ⑦ Variants of linear regression
  - Ridge regression
  - Polynomial regression
  - Lasso regression
  - Robust linear regression

1 Notations, vectors, matrices

2 Functions, derivative and gradient

3 Modeling of linear regression

- Deterministic perspective
- Probabilistic perspective

4 Learning of linear regression

- Analytical solution
- Gradient descent algorithm

5 Linear regression of multiple outputs

6 Linear regression for classification

7 Variants of linear regression

- Ridge regression
- Polynomial regression
- Lasso regression
- Robust linear regression

# Notations

- A **set** is an **unordered collection** of unique elements.
- We denote a set as a **calligraphic capital character**, for example,  $\mathcal{S}$ .
- A set of numbers can be **finite** (include a fixed amount of values). In this case, it is denoted using accolades, for example,  $\{1, 3, 18, 23, 235\}$  or  $\{x_1, x_2, x_3, \dots, x_d\}$ .
- A set can also be **infinite**.

# Notations

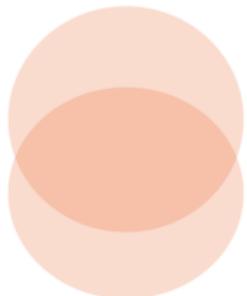
- A set can be **infinite** and include all values in some interval.
- If a set includes all values between  $a$  and  $b$ , **including  $a$  and  $b$** , it is denoted using brackets as  $[a, b]$ .
- If the set **does not include the values  $a$  and  $b$** , such a set is denoted using parentheses like this:  $(a, b)$ .
- For example, the set  $[0, 1]$  includes such values as 0, 0.0001, 0.25, 0.784, 0.9995, and 1.0.
- A special set denoted  $\mathcal{R}$  (or  $R, \mathbb{R}$ ) includes all **real numbers** from minus infinity to plus infinity.

# Notations

- **Intersection** of two sets:

$$\mathcal{S}_3 \leftarrow \mathcal{S}_1 \cap \mathcal{S}_2$$

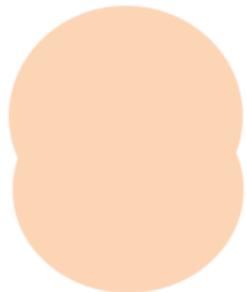
Example:  $\{1, 3, 5, 8\} \cap \{1, 8, 4\} = \{1, 8\}$



- **Union** of two sets:

$$\mathcal{S}_3 \leftarrow \mathcal{S}_1 \cup \mathcal{S}_2$$

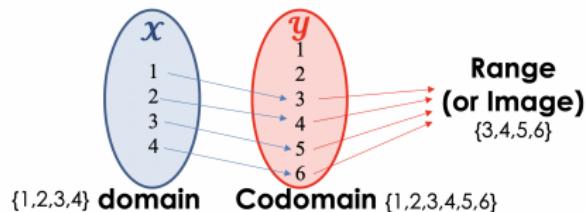
Example:  $\{1, 3, 5, 8\} \cup \{1, 8, 4\} = \{1, 3, 4, 5, 8\}$



- 1 Notations, vectors, matrices
- 2 Functions, derivative and gradient
- 3 Modeling of linear regression
  - Deterministic perspective
  - Probabilistic perspective
- 4 Learning of linear regression
  - Analytical solution
  - Gradient descent algorithm
- 5 Linear regression of multiple outputs
- 6 Linear regression for classification
- 7 Variants of linear regression
  - Ridge regression
  - Polynomial regression
  - Lasso regression
  - Robust linear regression

# Functions, derivative and gradient

- A **function** is a relation that associates each element  $x$  of a set  $\mathcal{X}$ , the **domain** of the function, to a single element  $y$  of another set  $\mathcal{Y}$ , the **codomain** of the function.
- A function usually has a name. If the function is called  $f$ , this relation is denoted  $y = f(x)$  (read  $f$  of  $x$ ), the element  $x$  is the **argument or input** of the function, and  $y$  is the **value of the function or the output**.
- The symbol that is used for representing the input is the variable of the function (we often say that  $f$  is a function of the variable  $x$ ).



# Functions, derivative and gradient

- A **scalar function** can also have vector argument such as,  $y = f(\mathbf{x})$ , or a scalar argument ( $y = f(\mathbf{x})$ ).
- A **vector function**, denoted as  $\mathbf{y} = \mathbf{f}(\mathbf{x})$ , is a function that returns  $\mathbf{y}$ , which can have either a vector argument ( $\mathbf{y} = \mathbf{f}(\mathbf{x})$ ) or a **scalar** argument ( $\mathbf{y} = \mathbf{f}(\mathbf{x})$ ).

# Functions, derivative and gradient

## Notation

- The notation  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  means that  $f$  is a function that maps real  $d$ -vectors to real numbers, i.e., it is a scalar-valued function of  $d$  dimension vectors.
- If  $\mathbf{x}$  is a  $d$ -vector, then  $f(\mathbf{x})$ , which is a scalar, denotes the value of the function  $f$  at  $\mathbf{x}$ . In the notation  $f(\mathbf{x})$ ,  $\mathbf{x}$  is referred to as the argument of the function

$$f(\mathbf{x}) = f(x_1, x_2, \dots, x_d)$$

# Functions, derivative and gradient

## Linear and Affine functions

- To describe a function  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  we have to specify its value for any possible argument  $\mathbf{x} \in \mathbb{R}^d$ .
- For example, we can define a function  $f : \mathbb{R}^4 \rightarrow \mathbb{R}$  by

$$f(\mathbf{x}) = x_1 + x_2 - 2x_3 - x_4$$

# Functions, derivative and gradient

## Linear and Affine functions

### Another example: The inner product function

- Suppose there is a  $d$ -vector. We can define a scalar-valued function  $f$  of  $d$ -vectors, given by

$$f(\mathbf{x}) = \mathbf{a}^\top \mathbf{x} = a_1 x_1 + a_2 x_2 + \dots + a_d x_d$$

for any  $d$ -vector  $\mathbf{x}$ .

- This function gives the inner product of its  $d$ -vector argument  $\mathbf{x}$  with some (fixed)  $d$ -vector  $\mathbf{a}$ .
- We can also think of  $f$  as forming a weighted sum of the elements of  $\mathbf{x}$ ; the elements of  $\mathbf{a}$  give the weights used in forming the weighted sum.

# Functions, derivative and gradient

## Linear and Affine functions

- A function  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  is **linear** if it satisfies the following **two properties**:
  - **Homogeneity**: For any  $d$ -vector  $\mathbf{x}$  and a scalar  $\alpha$ ,  $f(\alpha\mathbf{x}) = \alpha f(\mathbf{x})$ .
  - **Additivity**: For any  $d$ -vectors  $\mathbf{x}$  and  $\mathbf{y}$ ,  $f(\mathbf{x} + \mathbf{y}) = f(\mathbf{x}) + f(\mathbf{y})$

$$f(\mathbf{x}) = \mathbf{a}^\top \mathbf{x} = a_1 x_1 + a_2 x_2 + \dots + a_d x_d$$

for any  $d$ -vector  $\mathbf{x}$ .

- **Homogeneity** states that **scaling** the (vector) argument is the same as scaling the function value.
- **Additivity** says that **adding** (vector) arguments are the same as adding the function values.

# Functions, derivative and gradient

## Linear and Affine functions

### Superposition and linearity

- The inner product function  $f$  defined before satisfies the linearity property

$$\begin{aligned}f(\alpha \mathbf{x} + \beta \mathbf{y}) &= \mathbf{a}^\top (\alpha \mathbf{x} + \beta \mathbf{y}) \\&= \mathbf{a}^\top (\alpha \mathbf{x}) + \mathbf{a}^\top (\beta \mathbf{y}) \\&= \alpha (\mathbf{a}^\top \mathbf{x}) + \beta (\mathbf{a}^\top \mathbf{y}) \\&= \alpha f(\mathbf{x}) + \beta f(\mathbf{y})\end{aligned}$$

for all  $d$ -vectors  $\mathbf{x}, \mathbf{y}$ , and all scalars  $\alpha, \beta$ .

- This property is called **superposition** (which consists of homogeneity and additivity).
- A **function** that satisfies the superposition property is called **linear**

# Functions, derivative and gradient

## Linear and Affine functions

- If a function  $f$  is linear, superposition extends to linear combinations of any number of vectors:

$$f(\alpha_1 \mathbf{x}_1 + \dots + \alpha_k \mathbf{x}_k) = \alpha_1 f(\mathbf{x}_1) + \dots + \alpha_k f(\mathbf{x}_k)$$

for any  $d$ -vectors  $\mathbf{x}_1, \dots, \mathbf{x}_k$  and any scalars  $\alpha_1, \dots, \alpha_k$

# Functions, derivative and gradient

## Linear and Affine functions

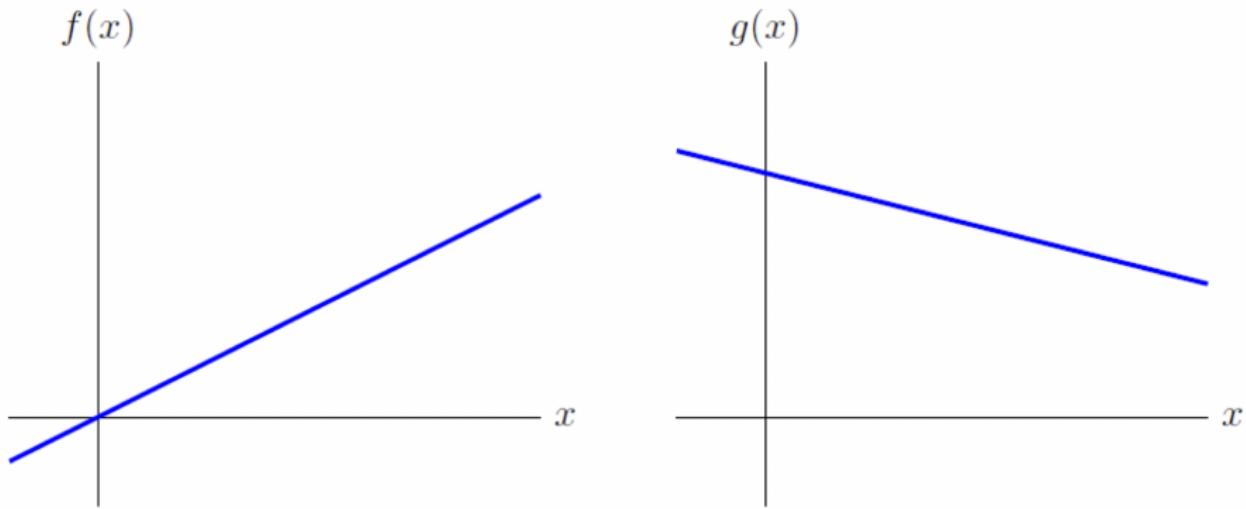
A function  $f : \mathbb{R}^d \rightarrow \mathbb{R}$  is **affine** if and only if it can be expressed as  $f(\mathbf{x}) = \mathbf{a}^\top \mathbf{x} + b$  for some  $d$ -vector  $\mathbf{a}$  and a scalar  $b$ , which is sometimes called the **offset**.

**Example:**

$$f(\mathbf{x}) = 2.3 - 2x_1 + 1.3x_2 - x_3$$

is affine, with  $b = 2.3$ ,  $\mathbf{a} = \begin{bmatrix} -2 \\ 1.3 \\ -1 \end{bmatrix}$ .

# Functions, derivative and gradient



**Figure 2.1** *Left.* The function  $f$  is linear. *Right.* The function  $g$  is affine, but not linear.

# Functions, derivative and gradient

- We say that  $f(x)$  has a **local minimum** at  $x = c$  if  $f(x) \geq f(c)$  for every  $x$  in some open interval around  $x = c$ .
- An **interval** is a set of real numbers with the property that any number that lies between two numbers in the set is also included in the set.
- An **open interval** does not include its endpoints and is denoted using parentheses. For example,  $(0, 1)$  means “all numbers greater than 0 and less than 1”
- The minimal value among all the **local minima** is called the **global minima**. See the illustration in the Figure on the next page.

# Functions, derivative and gradient

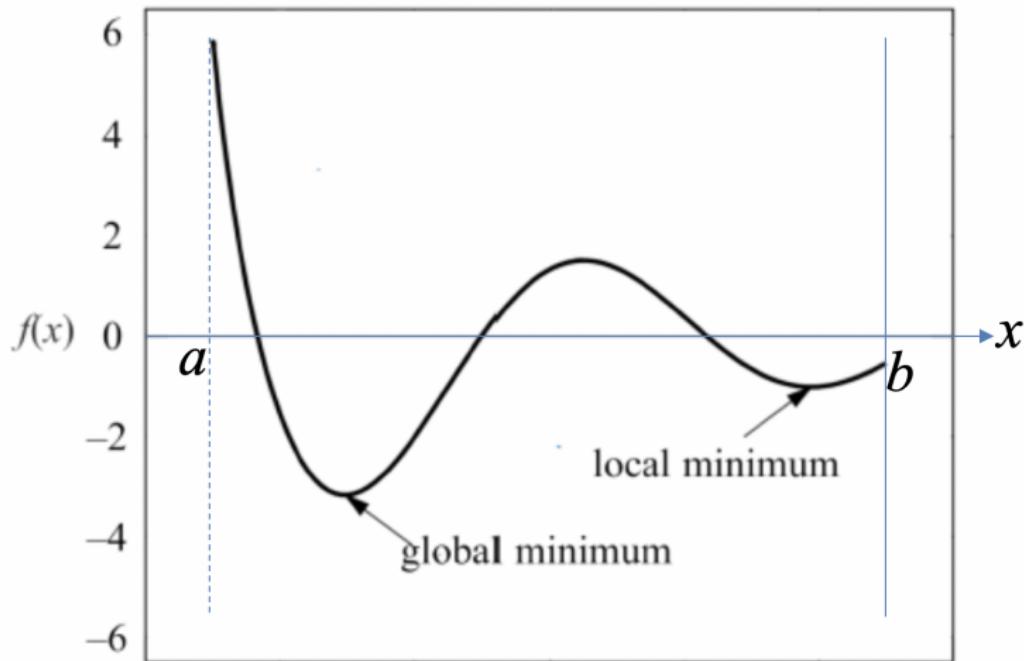


Figure: Local and global minima of a function.  $a < x < b$

# Functions, derivative and gradient

max vs arg max

- Given a set of values  $\mathcal{A} = \{a_1, a_2, \dots, a_m\}$ , the operator  $\max_{a \in \mathcal{A}} f(a)$  returns the highest value  $f(a)$  for all elements in the set  $\mathcal{A}$ .
- On the other hand, the operator  $\arg \max_{a \in \mathcal{A}} f(a)$  returns the element of the set  $\mathcal{A}$  that maximizes  $f(a)$ .
- Sometimes, when the set is implicit or explicit, we can write

$$\max_a f(a) \quad \text{or} \quad \arg \max_a f(a)$$

- Operator min and arg min operates in a similar manner.
- Note: **arg max** returns a value from the **domain** of the function and **max** returns from the **range (codomain)** of the function

# Functions, derivative and gradient

## Derivative and Gradient

- A **derivative**  $f'$  of a function  $f$  is a function or a value that describes how fast  $f$  grows (or decreases).
- If the derivative is a **constant** value, like 5 or  $-3$ , then the function grows (or decreases) constantly at any point  $x$  of its domain.
- If the derivative  $f'$  is **positive** at some point  $x$ , then the function  $f$  grows at this point.
- If the derivative  $f'$  is **negative** at some point  $x$ , then the function  $f$  decreases at this point.
- The **derivative of zero** at  $x$  means that the **function's slope** at  $x$  is **horizontal**.

# Functions, derivative and gradient

## Partial Derivative

- Differentiation of a **scalar** function *w.r.t.* a **vector**
- If  $f(\mathbf{w})$  is a **scalar function** of  $d$  variables,  $\mathbf{w}$  is a  $d \times 1$  vector, then differentiation of  $f(\mathbf{w})$  *w.r.t.*  $\mathbf{w}$  results in a  $d \times 1$  vector.

$$\frac{df(\mathbf{w})}{d\mathbf{w}} = \begin{bmatrix} \frac{\partial f}{\partial w_1} \\ \vdots \\ \frac{\partial f}{\partial w_d} \end{bmatrix}$$

This is referred to as the **gradient** of  $f(\mathbf{w})$  and written as  $\nabla_{\mathbf{w}} f$ .

# Functions, derivative and gradient

## Partial Derivative

- Differentiation of a **vector** function *w.r.t.* a **vector**
- If  $\mathbf{f}(\mathbf{w})$  is a vector function of size  $h \times 1$  and  $\mathbf{w}$  is a  $d \times 1$  vector, then differentiation of  $\mathbf{f}(\mathbf{w})$  *w.r.t.*  $\mathbf{w}$  results in a  $d \times h$  vector.

$$\frac{d\mathbf{f}(\mathbf{w})}{d\mathbf{w}} = \begin{bmatrix} \frac{\partial f_1}{\partial w_1} & \cdots & \frac{\partial f_h}{\partial w_1} \\ \vdots & \ddots & \vdots \\ \frac{\partial f_1}{\partial w_d} & \cdots & \frac{\partial f_h}{\partial w_d} \end{bmatrix}$$

- This is referred to as the **Jacobian** matrix of  $\mathbf{f}(\mathbf{w})$ , *i.e.*,

$$\mathbf{J} = \frac{d\mathbf{f}(\mathbf{w})}{d\mathbf{w}},$$
$$\mathbf{J}_{ij} = \frac{\partial f_j}{\partial w_i}.$$

# Functions, derivative and gradient

## Some Vector-Matrix Differentiation Formulations

$$\frac{d(\mathbf{X}^\top \mathbf{w})}{d\mathbf{w}} = \mathbf{X}, \text{ where } \mathbf{X} \text{ is not a function of } \mathbf{w}$$

$$\frac{d(\mathbf{y}^\top \mathbf{X}\mathbf{w})}{d\mathbf{w}} = \mathbf{X}^\top \mathbf{y}$$

$$\frac{d(\mathbf{w}^\top \mathbf{X}\mathbf{w})}{d\mathbf{w}} = (\mathbf{X} + \mathbf{X}^\top)\mathbf{w}$$

- Note that we adopt the **denominator layout** derivative. If you use the **numerator layout** derivative, then all the above results will be transposed.
- Both types are OK, but keep it **consistent** in all derivatives.
- Please refer to the following wiki page for more details:  
[https://en.wikipedia.org/wiki/Matrix\\_calculus](https://en.wikipedia.org/wiki/Matrix_calculus)

1 Notations, vectors, matrices

2 Functions, derivative and gradient

3 Modeling of linear regression

- Deterministic perspective
- Probabilistic perspective

4 Learning of linear regression

- Analytical solution
- Gradient descent algorithm

5 Linear regression of multiple outputs

6 Linear regression for classification

7 Variants of linear regression

- Ridge regression
- Polynomial regression
- Lasso regression
- Robust linear regression

# Linear regression

**Dataset:** A collection of  $m$  labeled examples  $\{(\mathbf{x}_i, y_i)\}_{i=1}^m$ , with  $\mathbf{x}_i \in \mathcal{X}$  being the  $d$ -dimensional feature vector of the  $i$ -th example, and  $y_i \in \mathcal{Y}$  being a real-valued target.

## Linear hypothesis function:

- We want to build a **linear** model, *i.e.*, a linear hypothesis function, as follows

$$f_{w_0, \dots, w_d}(\mathbf{x}) = w_0 + w_1 x_1 + w_2 x_2 + \cdots + w_d x_d,$$

where  $w_0$  is often called bias and  $w_1, \dots, w_d$  are often called coefficients.

- Let  $\mathbf{w} = [w_0, w_1, \dots, w_d]^\top$  and rewrite  $\mathbf{x} \leftarrow [1, x_1, \dots, x_d]^\top$  (augmented feature vector). Then the linear model is equivalent to

$$f_{\mathbf{w}}(\mathbf{x}) = \mathbf{w}^\top \mathbf{x}$$

- **Note:**  $f_{\mathbf{w}}$  is called **linear** due to the linearity *w.r.t.* the parameter vector  $\mathbf{w}$ , rather than *w.r.t.* the original feature vector  $[x_1, \dots, x_d]^\top$ .

## Task of linear regression:

- Using  $f_{\mathbf{w}}$  to approximate the ground-truth target function  $t : \mathcal{X} \rightarrow \mathcal{Y}$ .
- **Note:** If  $\mathcal{Y}$  is a **finite and discrete** set, then the task corresponds to a **classification problem**; if  $\mathcal{Y}$  is a **continuous** space, then the task corresponds to a **regression problem**.

# Linear regression: deterministic perspective

## Learning objective function

- To find the optimal values for  $\mathbf{w}$  which **minimizes** the following expression:

$$\frac{1}{m} \sum_{i=1}^m (f_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2$$

- In mathematics, the expression we minimize or maximize is called an **objective function**, or, simply, an **objective**.

# Linear regression: deterministic perspective

- The expression  $(f_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2$  in the above objective is called the **loss function**. It's a penalty measure of mismatch for example  $i$ .
- This particular choice of the loss function is called **squared error loss**.
- All model-based learning algorithms have a loss function and what we do to find the best model is we try to minimize the objective known as the **cost function**.
- In linear regression, the cost function is given by the average loss, also called the **empirical risk**.

# Linear regression: probabilistic perspective

- We assume that the relationship between the input variable/feature  $\mathbf{x}$  and the output variable  $y$  is

$$y = \mathbf{w}^\top \mathbf{x} + e, \text{ where } e \sim \mathcal{N}(0, \sigma^2), \quad (1)$$

where  $e$  is called **observation noise** or **residual error**, and it is independent with any specific input  $\mathbf{x}$ .

- Thus, the output  $y$  can also be seen as a random variable, and its conditional probability is formulated as

$$p(y|\mathbf{x}, \mathbf{w}) = \mathcal{N}(\mathbf{w}^\top \mathbf{x}, \sigma^2) \quad (2)$$

# Linear regression: probabilistic perspective

## Maximum log-likelihood estimation:

- The parameter  $\mathbf{w}$  can be learned by maximum log-likelihood estimation (MLE), given the training dataset  $D = \{(\mathbf{x}_i, y_i)\}_{i=1}^m$ , as follows

$$\mathbf{w}_{MLE} = \arg \max_{\mathbf{w}} \log \mathcal{L}(\mathbf{w}; D), \quad (3)$$

$$\begin{aligned} \log \mathcal{L}(\mathbf{w}; D) &= \log \left( \prod_{i=1}^m p(y_i | \mathbf{x}_i, \mathbf{w}) \right) = \sum_{i=1}^m \log \mathcal{N}(\mathbf{w}^\top \mathbf{x}_i, \sigma^2) \\ &= -m \log(\sigma(2\pi)^{\frac{1}{2}}) - \frac{1}{2\sigma^2} \sum_{i=1}^m (y_i - \mathbf{w}^\top \mathbf{x}_i)^2. \end{aligned} \quad (4)$$

- Removing the constants *w.r.t.*  $\mathbf{w}$ ,

$$\mathbf{w}_{MLE} = \arg \min_{\mathbf{w}} \frac{1}{2} \sum_{i=1}^m (y_i - \mathbf{w}^\top \mathbf{x}_i)^2, \quad (5)$$

which is exactly the same with the cost function from the deterministic perspective.

- 1 Notations, vectors, matrices
- 2 Functions, derivative and gradient
- 3 Modeling of linear regression
  - Deterministic perspective
  - Probabilistic perspective
- 4 Learning of linear regression
  - Analytical solution
  - Gradient descent algorithm
- 5 Linear regression of multiple outputs
- 6 Linear regression for classification
- 7 Variants of linear regression
  - Ridge regression
  - Polynomial regression
  - Lasso regression
  - Robust linear regression

# Linear regression with analytical solution

## Learning (Training)

- Consider the set of feature vector  $\mathbf{x}_i$  and target output  $y_i$  indexed by  $i = 1, \dots, m$ , then a linear model  $f_{\mathbf{w}}(\mathbf{x}) = \mathbf{x}^\top \mathbf{w} = \sum_{j=0}^d w_j x_j$  can be packed as

$$f_{\mathbf{w}}(\mathbf{X}) \Leftrightarrow \mathbf{y} = \begin{bmatrix} y_1 \\ \vdots \\ y_m \end{bmatrix}$$

Learning model                      Learning target vector

$$= \begin{bmatrix} \mathbf{x}_1^\top \mathbf{w} \\ \vdots \\ \mathbf{x}_m^\top \mathbf{w} \end{bmatrix}$$

$$\text{where } \mathbf{x}_i^\top \mathbf{w} = [1, x_1, \dots, x_d]_i \quad \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_d \end{bmatrix} \text{ and } \mathbf{X} = \begin{bmatrix} \mathbf{x}_1^\top \\ \vdots \\ \mathbf{x}_m^\top \end{bmatrix}$$

**Note:** The bias term is responsible for shifting the line/plane up or down.

# Linear Regression

## Least Squares Regression

- In vector-matrix notation, the squared error function for all the  $m$  samples can be written compactly using  $\mathbf{e} = \mathbf{X}\mathbf{w} - \mathbf{y}$  :

$$\begin{aligned} J(\mathbf{w}) &= \sum_{i=1}^m (f_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2 \\ &= \sum_{i=1}^m e_i^2 \\ &= \mathbf{e}^\top \mathbf{e} \\ &= (\mathbf{X}\mathbf{w} - \mathbf{y})^\top (\mathbf{X}\mathbf{w} - \mathbf{y}) \\ &= (\mathbf{w}^\top \mathbf{X}^\top - \mathbf{y}^\top) (\mathbf{X}\mathbf{w} - \mathbf{y}) \\ &= \mathbf{w}^\top \mathbf{X}^\top \mathbf{X}\mathbf{w} - \mathbf{w}^\top \mathbf{X}^\top \mathbf{y} - \mathbf{y}^\top \mathbf{X}\mathbf{w} + \mathbf{y}^\top \mathbf{y} \\ &= \mathbf{w}^\top \mathbf{X}^\top \mathbf{X}\mathbf{w} - 2\mathbf{y}^\top \mathbf{X}\mathbf{w} + \mathbf{y}^\top \mathbf{y} \end{aligned}$$

# Linear Regression

Differentiating  $J(\mathbf{w})$  with respect to  $\mathbf{w}$  and setting the result to 0 :

$$\begin{aligned}\frac{\partial}{\partial \mathbf{w}} J(\mathbf{w}) &= \mathbf{0} \\ \frac{\partial}{\partial \mathbf{w}} (\mathbf{w}^\top \mathbf{X}^\top \mathbf{X} \mathbf{w} - 2\mathbf{y}^\top \mathbf{X} \mathbf{w} + \mathbf{y}^\top \mathbf{y}) &= \mathbf{0} \\ \Rightarrow 2\mathbf{X}^\top \mathbf{X} \mathbf{w} - 2\mathbf{X}^\top \mathbf{y} &= \mathbf{0} \\ \Rightarrow 2\mathbf{X}^\top \mathbf{X} \mathbf{w} &= 2\mathbf{X}^\top \mathbf{y}\end{aligned}$$

If  $\mathbf{X}^\top \mathbf{X}$  is invertible, then

$$\begin{aligned}\text{Learning: } \hat{\mathbf{w}} &= (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y} \\ \text{Prediction: } f_{\mathbf{w}}(\mathbf{X}_{\text{new}}) &= \mathbf{X}_{\text{new}} \hat{\mathbf{w}}\end{aligned}$$

# Linear Regression

## Example 1

$$\begin{bmatrix} 1 & -9 \\ 1 & -7 \\ 1 & -5 \\ 1 & 1 \\ 1 & 5 \\ 1 & 9 \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \end{bmatrix} = \begin{bmatrix} -6 \\ -6 \\ -4 \\ -1 \\ 1 \\ 4 \end{bmatrix}$$

$\mathbf{X} \quad \mathbf{w} \quad \mathbf{y}$

$$\begin{aligned} \{(x_i, y_i)\}_{i=1}^m \\ \{x = -9\} \rightarrow \{y = -6\} \\ \{x = -7\} \rightarrow \{y = -6\} \\ \{x = -5\} \rightarrow \{y = -4\} \\ \{x = 1\} \rightarrow \{y = -1\} \\ \{x = 5\} \rightarrow \{y = 1\} \\ \{x = 9\} \rightarrow \{y = 4\} \end{aligned}$$

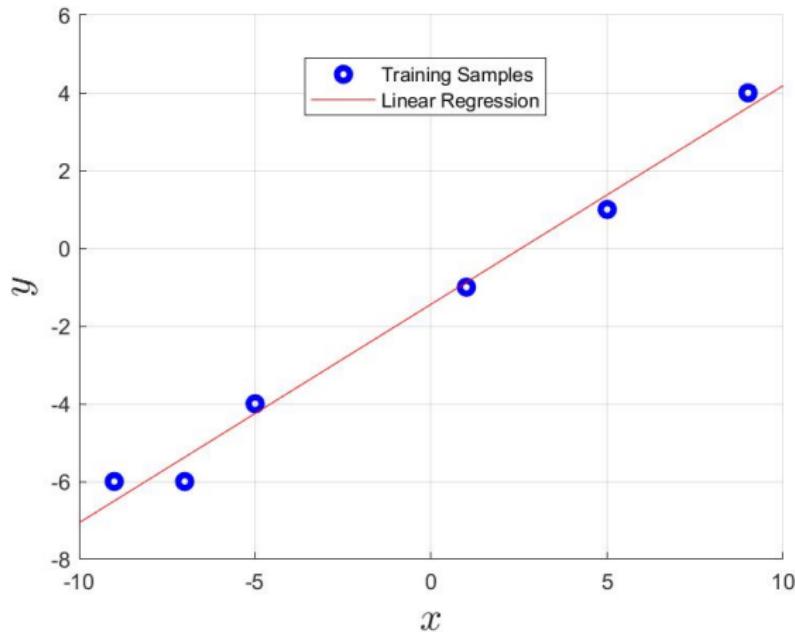
This set of linear equations has NO exact solution.

$$\hat{\mathbf{w}} = \mathbf{X}^\dagger \mathbf{y} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

However ( $\mathbf{X}^T \mathbf{X}$  is invertible)

$$= \begin{bmatrix} 6 & -6 \\ -6 & 262 \end{bmatrix}^{-1} \begin{bmatrix} 1 & 1 & 1 & 1 & 1 & 1 \\ -9 & -7 & -5 & 1 & 5 & 9 \end{bmatrix} \begin{bmatrix} -6 \\ -6 \\ -4 \\ -1 \\ 1 \\ 4 \end{bmatrix} = \begin{bmatrix} -1.4375 \\ 0.5625 \end{bmatrix}$$

# Linear regression



$$\begin{aligned}\hat{y} &= \mathbf{X}\hat{\mathbf{w}} \\ &= \mathbf{X} \begin{bmatrix} -1.4375 \\ 0.5625 \end{bmatrix} \\ y &= -1.4375 + 0.5625x\end{aligned}$$

Prediction:

$$\begin{aligned}\{x = -1\} \rightarrow \{y = ?\} \\ \hat{y} &= [1 \ 1] \begin{bmatrix} -1.4375 \\ 0.5625 \end{bmatrix} \\ &= -2\end{aligned}$$

Linear Regression for one-dimensional examples.

# Linear Regression

## Example 2

$$\{(\mathbf{x}_i, y_i)\}_{i=1}^m \quad \begin{aligned} \{x_1 = 1, x_2 = 1, x_3 = 1\} &\rightarrow \{y = 1\} \\ \{x_1 = 1, x_2 = -1, x_3 = 1\} &\rightarrow \{y = 0\} \\ \{x_1 = 1, x_2 = 1, x_3 = 3\} &\rightarrow \{y = 2\} \\ \{x_1 = 1, x_2 = 1, x_3 = 0\} &\rightarrow \{y = -1\} \end{aligned}$$

$$\mathbf{X} = \begin{bmatrix} 1 & 1 & 1 \\ 1 & -1 & 1 \\ 1 & 1 & 3 \\ 1 & 1 & 0 \end{bmatrix}, \mathbf{w} = \begin{bmatrix} w_1 \\ w_2 \\ w_3 \end{bmatrix}, \mathbf{y} = \begin{bmatrix} 1 \\ 0 \\ 2 \\ -1 \end{bmatrix} \quad \mathbf{w}^T \mathbf{X}^T$$

This set of linear equations has NO exact solution.

$$\hat{\mathbf{w}} = \mathbf{X}^\dagger \mathbf{y} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y} \quad \text{However } (\mathbf{X}^T \mathbf{X} \text{ is invertible)}$$

$$= \begin{bmatrix} 4 & 2 & 5 \\ 2 & 4 & 3 \\ 5 & 3 & 11 \end{bmatrix}^{-1} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & -1 & 1 & 1 \\ 1 & 1 & 3 & 0 \end{bmatrix} \begin{bmatrix} 1 \\ 0 \\ 2 \\ -1 \end{bmatrix} \approx \begin{bmatrix} -0.75 \\ 0.18 \\ 0.93 \end{bmatrix} \quad \begin{array}{l} \text{(Least squares} \\ \text{approximation).} \end{array}$$

# Linear Regression

**Prediction:**

$$\{x_1 = 1, x_2 = 6, x_3 = 8\} \rightarrow \{y = ?\}$$

$$\{x_1 = 1, x_2 = 0, x_3 = -1\} \rightarrow \{y = ?\}$$

$$\begin{aligned}\hat{y} &= \begin{bmatrix} 1 & 6 & 8 \\ 1 & 0 & -1 \end{bmatrix} \begin{bmatrix} -0.75 \\ 0.18 \\ 0.93 \end{bmatrix} \\ &= \begin{bmatrix} 7.7500 \\ -1.6786 \end{bmatrix}\end{aligned}$$

# Linear regression solved by gradient descent

- The linear regression is formulated for the following optimization problem

$$\mathbf{w}^* = \arg \min_{\mathbf{w}} J(\mathbf{w}), \quad J(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^m (\mathbf{x}_i^\top \mathbf{w} - y_i)^2 = \frac{1}{2} \|\mathbf{X}\mathbf{w} - \mathbf{y}\|^2 \quad (6)$$

- $\mathbf{w}$  can be updated by [gradient descent algorithm](#),

$$\mathbf{w} \leftarrow \mathbf{w} - \alpha \frac{\partial J(\mathbf{w})}{\partial \mathbf{w}}, \quad \frac{\partial J(\mathbf{w})}{\partial \mathbf{w}} = \mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{y}) \quad (7)$$

where  $\alpha$  is called step size or learning rate.

- Does gradient descent always converge to the optimal solution? ([Plot the update trajectory of gradient descent on loss curve](#))

# Closed-form solution vs. gradient descent

| Closed-form solution<br>$\mathbf{w} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$ | Gradient descent<br>$\mathbf{w} \leftarrow \mathbf{w} - \alpha \mathbf{X}^\top (\mathbf{X}\mathbf{w} - \mathbf{y})$ |
|---|---|
| No hyper-parameter  | Needs to choose $\alpha$  |
| No need to iterate  | Needs many iterations   |
| Complexity $O(d^3 + md^2)$  | Complexity $O(T \times md)$   |
| Slow if $d$ is very large   | Works well when $d$ is large  |

Thus, you can choose between the above two solutions according to the dimensionality of your training data:

- When the training data is very high-dimensional, *i.e.*,  $d$  is very **large**, then it is better to choose **gradient descent algorithm**
- When the training data is not high-dimensional, *i.e.*,  $d$  is very **small**, then it is better to choose **closed-form solution**

- 1 Notations, vectors, matrices
- 2 Functions, derivative and gradient
- 3 Modeling of linear regression
  - Deterministic perspective
  - Probabilistic perspective
- 4 Learning of linear regression
  - Analytical solution
  - Gradient descent algorithm
- 5 Linear regression of multiple outputs
- 6 Linear regression for classification
- 7 Variants of linear regression
  - Ridge regression
  - Polynomial regression
  - Lasso regression
  - Robust linear regression

# Linear regression with single output

When considering the entire set of data indexed by  $i = 1, \dots, m$ , a linear model  $f_{\mathbf{w}}(\mathbf{x}) = \mathbf{x}^\top \mathbf{w} = \sum_{j=0}^d w_j x_j$ , where  $x_0 \equiv 1$ , can be packed as

$$f_{\mathbf{w}}(\mathbf{X}) = \mathbf{X}\mathbf{w} \leftarrow \text{Vector function}$$

$$= \begin{bmatrix} \mathbf{x}_1^\top \mathbf{w} \\ \vdots \\ \mathbf{x}_m^\top \mathbf{w} \end{bmatrix} \quad \text{where} \quad \mathbf{x}_i^\top \mathbf{w} = [1, x_{i,1}, \dots, x_{i,d}] \begin{bmatrix} w_0 \\ w_1 \\ \vdots \\ w_d \end{bmatrix} \quad (8)$$

Primal solution:  $\hat{\mathbf{w}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$

**Note:** The **bias term** is responsible for **shifting the line/plane** up or down.

# Linear regression with multiple outputs

- When considering the entire set of data indexed by  $i = 1, \dots, m$ , a linear model  $\mathbf{f}_W(\mathbf{x}) = \mathbf{x}^\top \mathbf{W}$  can be packed as

$$\mathbf{f}_W(\mathbf{X}) = \mathbf{X}\mathbf{W} \leftarrow \text{Matrix function}$$

$$\begin{array}{c} \text{Sample 1} \xrightarrow{\quad} \mathbf{x}_1^T \\ \vdots \\ \text{Sample } m \xrightarrow{\quad} \mathbf{x}_m^T \end{array} = \begin{bmatrix} \mathbf{x}_1^T \\ \vdots \\ \mathbf{x}_m^T \end{bmatrix} \mathbf{W} = \begin{bmatrix} 1 & x_{1,1} & \dots & x_{1,d} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & x_{m,1} & \dots & x_{m,d} \end{bmatrix} \begin{bmatrix} w_{0,1} & \dots & w_{0,h} \\ w_{1,1} & \dots & w_{1,h} \\ \vdots & \ddots & \vdots \\ w_{d,1} & \dots & w_{d,h} \end{bmatrix}$$
  
$$\begin{array}{c} \text{Sample 1's output} \xrightarrow{\quad} f_{1,1} & \dots & f_{1,h} \\ \vdots \\ \text{Sample } m \text{'s output} \xrightarrow{\quad} f_{m,1} & \dots & f_{m,h} \end{array} = \begin{bmatrix} f_{1,1} & \dots & f_{1,h} \\ \vdots & \ddots & \vdots \\ f_{m,1} & \dots & f_{m,h} \end{bmatrix} \underbrace{\qquad\qquad\qquad}_{h} \Big\} m$$

# Linear regression with multiple outputs

In matrix-matrix notation, the squared error loss function can be written compactly using  $\mathbf{E} = \mathbf{X}\mathbf{W} - \mathbf{Y}$  :

$$\begin{aligned} J(\mathbf{W}) &= \text{trace} (\mathbf{E}^\top \mathbf{E}) \\ &= \text{trace} [(\mathbf{X}\mathbf{W} - \mathbf{Y})^\top (\mathbf{X}\mathbf{W} - \mathbf{Y})] \end{aligned}$$

If  $\mathbf{X}^\top \mathbf{X}$  is invertible, then

$$\begin{aligned} \text{Learning: } \widehat{\mathbf{W}} &= (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{Y} \\ \text{Prediction: } \mathbf{f}_{\mathbf{W}} (\mathbf{X}_{\text{new}}) &= \mathbf{X}_{\text{new}} \widehat{\mathbf{W}} \end{aligned}$$

Assumption: the error terms  $\mathbf{e}_k^\top \mathbf{e}_k$  are independent for all  $k = 1, \dots, h$

## Linear regression with multiple outputs

$$J(\mathbf{W}) = \text{trace}(\mathbf{E}^T \mathbf{E})$$

$$= \text{trace}\left(\begin{bmatrix} \mathbf{e}_1^T \\ \vdots \\ \mathbf{e}_h^T \end{bmatrix} [\mathbf{e}_1 \ \ \mathbf{e}_2 \ \ \dots \ \ \mathbf{e}_h]\right)$$

$$= \text{trace}\left(\begin{bmatrix} \mathbf{e}_1^T \mathbf{e}_1 & \mathbf{e}_1^T \mathbf{e}_2 & \dots & \mathbf{e}_1^T \mathbf{e}_h \\ \mathbf{e}_2^T \mathbf{e}_1 & \mathbf{e}_2^T \mathbf{e}_2 & \dots & \mathbf{e}_2^T \mathbf{e}_h \\ \vdots & \vdots & \ddots & \vdots \\ \mathbf{e}_h^T \mathbf{e}_1 & \mathbf{e}_h^T \mathbf{e}_2 & \dots & \mathbf{e}_h^T \mathbf{e}_h \end{bmatrix}\right) = \sum_{k=1}^h \mathbf{e}_k^T \mathbf{e}_k$$

# Linear regression

$$\{\mathbf{x}_i, \mathbf{y}_i\}_{i=1}^m \quad \{x_1 = 1, x_2 = 1, x_3 = 1\} \rightarrow \{y_1 = 1, y_2 = 0\}$$

Example

$$\{x_1 = 1, x_2 = -1, x_3 = 1\} \rightarrow \{y_1 = 0, y_2 = 1\}$$

$$\{x_1 = 1, x_2 = 1, x_3 = 3\} \rightarrow \{y_1 = 2, y_2 = -1\}$$

$$\{x_1 = 1, x_2 = 1, x_3 = 0\} \rightarrow \{y_1 = -1, y_2 = 3\}$$

$$\begin{bmatrix} 1 & 1 & 1 \\ 1 & -1 & 1 \\ 1 & 1 & 3 \\ 1 & 1 & 0 \end{bmatrix} \begin{bmatrix} w_{1,1} & w_{1,2} \\ w_{2,1} & w_{2,2} \\ w_{3,1} & w_{3,2} \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 0 & 1 \\ 2 & -1 \\ -1 & 3 \end{bmatrix}$$

$\mathbf{X} \qquad \mathbf{W} \qquad \mathbf{Y}$

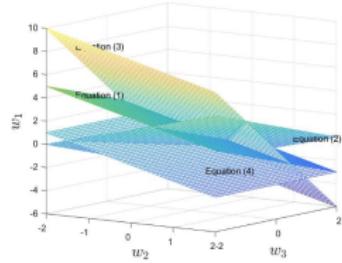
This set of linear equations has NO exact solution.

$$\hat{\mathbf{W}} = \mathbf{X}^\dagger \mathbf{Y} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y} \quad \text{However } (\mathbf{X}^T \mathbf{X} \text{ is invertible})$$

$$= \begin{bmatrix} 4 & 2 & 5 \\ 2 & 4 & 3 \\ 5 & 3 & 11 \end{bmatrix}^{-1} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & -1 & 1 & 1 \\ 1 & 1 & 3 & 0 \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & 1 \\ 2 & -1 \\ -1 & 3 \end{bmatrix} = \begin{bmatrix} -0.75 & 2.25 \\ 0.1786 & 0.0357 \\ 0.9286 & -1.2143 \end{bmatrix} \quad \text{(Least squares approximation).}$$

# Linear regression with multiple outputs

Example:



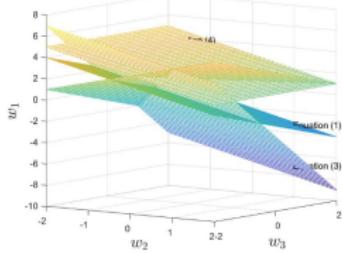
Prediction:

$$\{x_1 = 1, x_2 = 6, x_3 = 8\} \rightarrow \{y_1 = ?, y_2 = ?\}$$
$$\{x_1 = 1, x_2 = 0, x_3 = -1\} \rightarrow \{y_1 = ?, y_2 = ?\}$$

$$\hat{\mathbf{Y}} = \mathbf{X}_t \hat{\mathbf{W}}$$

$$= \begin{bmatrix} 1 & 6 & 8 \\ 1 & 0 & -1 \end{bmatrix} \begin{bmatrix} -0.75 & 2.25 \\ 0.1786 & 0.0357 \\ 0.9286 & -1.2143 \end{bmatrix}$$

$$= \begin{bmatrix} 7.75 & -7.25 \\ -1.6786 & 3.4643 \end{bmatrix}$$



- 1 Notations, vectors, matrices
- 2 Functions, derivative and gradient
- 3 Modeling of linear regression
  - Deterministic perspective
  - Probabilistic perspective
- 4 Learning of linear regression
  - Analytical solution
  - Gradient descent algorithm
- 5 Linear regression of multiple outputs
- 6 Linear regression for classification
- 7 Variants of linear regression
  - Ridge regression
  - Polynomial regression
  - Lasso regression
  - Robust linear regression

# Linear regression for classification

## Dataset:

- We have a collection of  $m$  labeled examples  $\{(\mathbf{x}_i, y_i)\}_{i=1}^m$ , with  $\mathbf{x}_i \in \mathcal{X}$  being the  $d$ -dimensional feature vector of the  $i$ -th example, and  $y_i \in \mathcal{Y}$  being a real-valued target.

## Linear hypothesis function:

- We want to build a linear model  $f_{\mathbf{w}}(\mathbf{x})$ , i.e., linear hypothesis function,

$$f_{\mathbf{w}}(\mathbf{x}) = \mathbf{x}^\top \mathbf{w},$$

where  $\mathbf{w} = [w_0, w_1, \dots, w_d]^\top$  is a  $(d+1)$ -dimensional vector of parameters,  $w_0$  is the bias parameter, and  $\mathbf{x}$  is the  $(d+1)$ -dimensional augmented feature vector, i.e.,  $\mathbf{x} = [1, x_1, \dots, x_d]^\top$ .

## Task of linear regression:

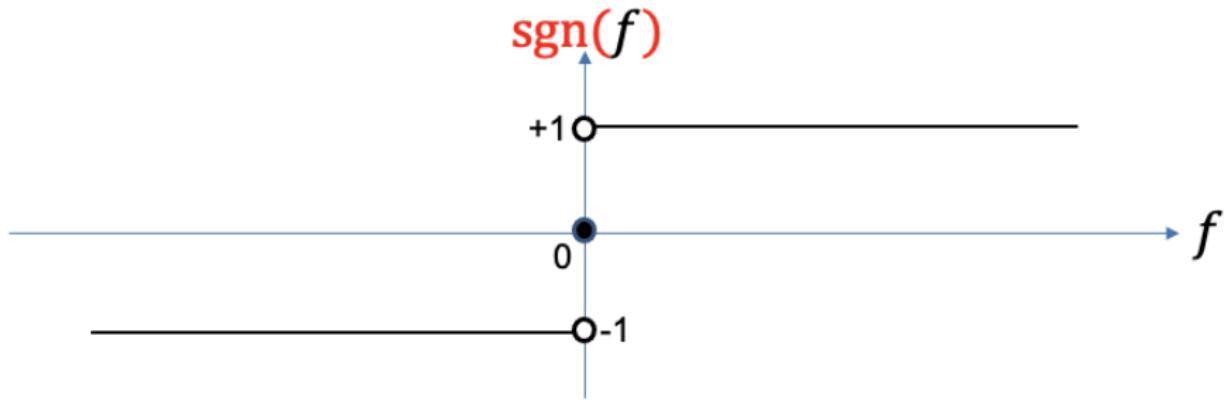
- Using the linear model  $f_{\mathbf{w}}$  to approximate the ground-truth target function  $t : \mathcal{X} \rightarrow \mathcal{Y}$ .
- Note: If  $\mathcal{Y}$  is a **finite and discrete** set, then the task corresponds to a **classification problem**; if  $\mathcal{Y}$  is a **continuous** space, then the task corresponds to a regression problem.

# Linear regression for classification

**Binary Classification:** If  $\mathbf{X}^\top \mathbf{X}$  is invertible, then

Learning:  $\hat{\mathbf{w}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}, \quad y_i \in \{-1, +1\}, i = 1, \dots, m$

Prediction:  $f_{\mathbf{w}}(\mathbf{x}_{new}) = \text{sgn}(\mathbf{x}_{new}^\top \hat{\mathbf{w}})$  ( for each row  $\mathbf{x}_{new}^\top$  of  $\mathbf{X}_{new}$  )



# Linear regression for classification

## Example

$\{\mathbf{x}_i, y_i\}_{i=1}^m$

- $\{x = -9\} \rightarrow \{y = -1\}$
- $\{x = -7\} \rightarrow \{y = -1\}$
- $\{x = -5\} \rightarrow \{y = -1\}$
- $\{x = -1\} \rightarrow \{y = +1\}$
- $\{x = 5\} \rightarrow \{y = +1\}$
- $\{x = 9\} \rightarrow \{y = +1\}$

$$\begin{bmatrix} 1 & -9 \\ 1 & -7 \\ 1 & -5 \\ 1 & 1 \\ 1 & 5 \\ 1 & 9 \end{bmatrix} \begin{bmatrix} w_0 \\ w_1 \end{bmatrix} = \begin{bmatrix} -1 \\ -1 \\ -1 \\ 1 \\ 1 \\ 1 \end{bmatrix}$$

$\mathbf{X}$        $\mathbf{w}$        $\mathbf{y}$

This set of linear equations has NO exact solution.

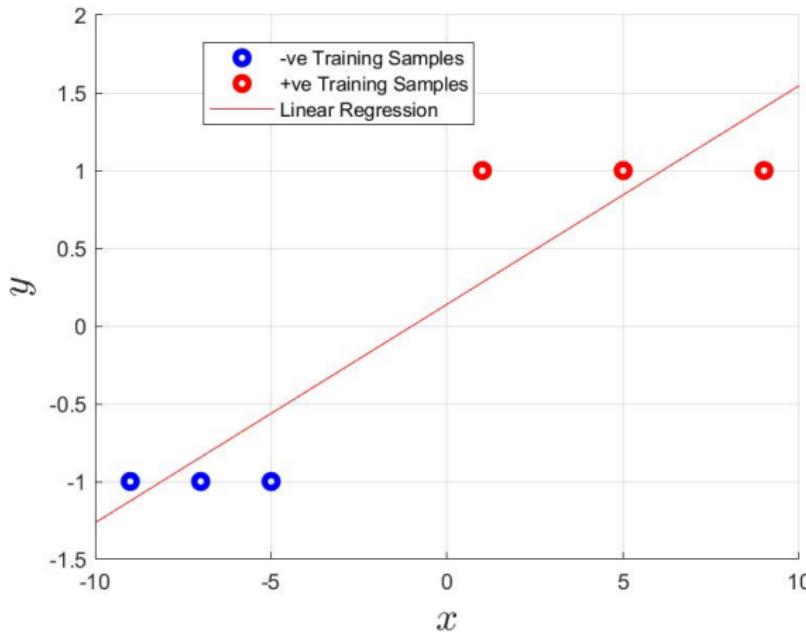
However  
 $(\mathbf{X}^T \mathbf{X})$  is invertible

$$\hat{\mathbf{w}} = \mathbf{X}^\dagger \mathbf{y} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{y}$$

$$= \begin{bmatrix} 6 & -6 \\ -6 & 262 \end{bmatrix}^{-1} \begin{bmatrix} 1 & 1 & 1 & 1 & 1 \\ -9 & -7 & -5 & 1 & 5 \\ 1 & 1 & 1 & 1 & 1 \end{bmatrix} \begin{bmatrix} -1 \\ -1 \\ -1 \\ 1 \\ 1 \\ 1 \end{bmatrix} = \begin{bmatrix} 0.1406 \\ 0.1406 \end{bmatrix}$$

(Least squares approximation).

# Linear regression for classification



$$\begin{aligned}\hat{y} &= \text{sgn}(\mathbf{X}\hat{\mathbf{w}}) \\ &= \text{sgn}\left(\mathbf{X} \begin{bmatrix} 0.1406 \\ 0.1406 \end{bmatrix}\right)\end{aligned}$$

$$y = 0.1406 + 0.1406x$$

Prediction:

$$\{x = -2\} \rightarrow \{y = ?\}$$

$$\begin{aligned}\hat{y} &= \text{sgn}([1 \ -2] \begin{bmatrix} 0.1406 \\ 0.1406 \end{bmatrix}) \\ &= \text{sgn}(-0.1406) = -1\end{aligned}$$

Linear regression for one-dimensional classification.

# Linear regression for classification

## Linear Methods for Multi-Category Classification:

If  $\mathbf{X}^\top \mathbf{X}$  is invertible, then

Learning:  $\widehat{\mathbf{W}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{Y}, \quad \mathbf{Y} \in \mathbf{R}^{m \times C}$

Prediction:  $f_{\mathbf{w}}(\mathbf{x}_{\text{new}}) = \arg \max_{i=1, \dots, C} (\mathbf{x}_{\text{new}}^\top \widehat{\mathbf{W}})$  (for each row  $\mathbf{x}_{\text{new}}^\top$  of  $\mathbf{X}_{\text{new}}$ )

Each row (of  $i=1, \dots, m$ ) in  $\mathbf{Y}$  has a one-hot assignment:

e.g., target for class-1 is labelled as  $\mathbf{y}_i^\top = [1, 0, 0, \dots, 0]$  for the  $i$  th sample,

target for class-2 is labelled as  $\mathbf{y}_i^\top = [0, 1, 0, \dots, 0]$  for the  $i$  th sample,

target for class-C is labelled as  $\mathbf{y}_i^\top = [0, 0, \dots, 0, 1]$  for the  $i$  th sample.

# Linear regression for classification

## Example

$$\{\mathbf{x}_i, \mathbf{y}_i\}_{i=1}^m \quad \begin{aligned} \{x_1 = 1, x_2 = 1, x_3 = 1\} &\rightarrow \{y_1 = 1, y_2 = 0, y_3 = 0\} \\ \{x_1 = 1, x_2 = -1, x_3 = 1\} &\rightarrow \{y_1 = 0, y_2 = 1, y_3 = 0\} \\ \{x_1 = 1, x_2 = 1, x_3 = 3\} &\rightarrow \{y_1 = 1, y_2 = 0, y_3 = 0\} \\ \{x_1 = 1, x_2 = 1, x_3 = 0\} &\rightarrow \{y_1 = 0, y_2 = 0, y_3 = 1\} \end{aligned}$$

$$\begin{bmatrix} 1 & 1 & 1 \\ 1 & -1 & 1 \\ 1 & 1 & 3 \\ 1 & 1 & 0 \end{bmatrix} \begin{bmatrix} w_{1,1} & w_{1,2} & w_{1,3} \\ w_{2,1} & w_{2,2} & w_{2,3} \\ w_{3,1} & w_{3,2} & w_{3,3} \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

$\mathbf{X} \qquad \mathbf{W} \qquad \mathbf{Y}$

This set of linear equations has NO exact solution.

$$\widehat{\mathbf{W}} = \mathbf{X}^\dagger \mathbf{Y} = (\mathbf{X}^T \mathbf{X})^{-1} \mathbf{X}^T \mathbf{Y} \quad \text{However } (\mathbf{X}^T \mathbf{X} \text{ is invertible)} \quad \text{(Least squares approximation)}$$

$$= \begin{bmatrix} 4 & 2 & 5 \\ 2 & 4 & 3 \\ 5 & 3 & 11 \end{bmatrix}^{-1} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & -1 & 1 & 1 \\ 1 & 1 & 3 & 0 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 1 & 0 & 0 \\ 0 & 0 & 1 \end{bmatrix} = \begin{bmatrix} 0 & 0.5 & 0.5 \\ 0.2857 & -0.5 & 0.2143 \\ 0.2857 & 0 & -0.2857 \end{bmatrix}$$

# Linear regression for classification

## Example

Prediction:

$$\{x_1 = 1, x_2 = 6, x_3 = 8\} \rightarrow \{ \text{class 1, 2, or 3 ? } \}$$

$$\{x_1 = 1, x_2 = 0, x_3 = -1\} \rightarrow \{ \text{class 1, 2, or 3 ? } \}$$

$$\widehat{\mathbf{Y}} = \mathbf{X}_t \widehat{\mathbf{W}} = \arg \max_{i=1, \dots, C} \left( \begin{bmatrix} 1 & 6 & 8 \\ 1 & 0 & -1 \end{bmatrix} \begin{bmatrix} 0 & 0.5 & 0.5 \\ 0.2857 & -0.5 & 0.2143 \\ 0.2857 & 0 & -0.2857 \end{bmatrix} \right)$$

$$= \arg \max_{i=1, \dots, C} \left( \begin{bmatrix} 4 & -2.50 & -0.50 \\ -0.2587 & 0.50 & 0.7857 \end{bmatrix} \right)$$

Position of the largest number determines the class label

$$= \begin{bmatrix} 1 \\ 3 \end{bmatrix} \xrightarrow{\text{Class-1}} \xrightarrow{\text{Class-3}}$$

- 1 Notations, vectors, matrices
- 2 Functions, derivative and gradient
- 3 Modeling of linear regression
  - Deterministic perspective
  - Probabilistic perspective
- 4 Learning of linear regression
  - Analytical solution
  - Gradient descent algorithm
- 5 Linear regression of multiple outputs
- 6 Linear regression for classification
- 7 Variants of linear regression
  - Ridge regression
  - Polynomial regression
  - Lasso regression
  - Robust linear regression

# Ridge regression

## Motivation 1:

Recall the learning computation:  $\hat{\mathbf{w}} = (\mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}$ .

We cannot guarantee that the matrix  $\mathbf{X}^\top \mathbf{X}$  is invertible.

# Ridge regression

Here on, we shall focus on single output  $y$  in derivations in the sequel.

$$\min_{\mathbf{w}} \sum_{i=1}^m (f_{\mathbf{w}}(\mathbf{x}_i) - y_i)^2 + \lambda \tilde{\mathbf{w}}^\top \tilde{\mathbf{w}}, \text{ where } \tilde{\mathbf{w}} = \hat{\mathbf{I}}_d \mathbf{w} = [0, w_1, w_2, \dots, w_d]^\top,$$

$\hat{\mathbf{I}}_d \in \mathbb{R}^{(d+1) \times (d+1)}$  is defined by setting the  $(1, 1)$  entry in the  $d + 1$  dimensional identity matrix  $\hat{\mathbf{I}}_{d+1}$  as 0. Note: The bias/offset  $w_0$  is NOT included in the  $\ell_2$  regularization term, as it just affects the function's height, not its complexity.

**Linear Model:**  $\min_{\mathbf{w}} (\mathbf{X}\mathbf{w} - \mathbf{y})^\top (\mathbf{X}\mathbf{w} - \mathbf{y}) + \lambda \tilde{\mathbf{w}}^\top \tilde{\mathbf{w}}$

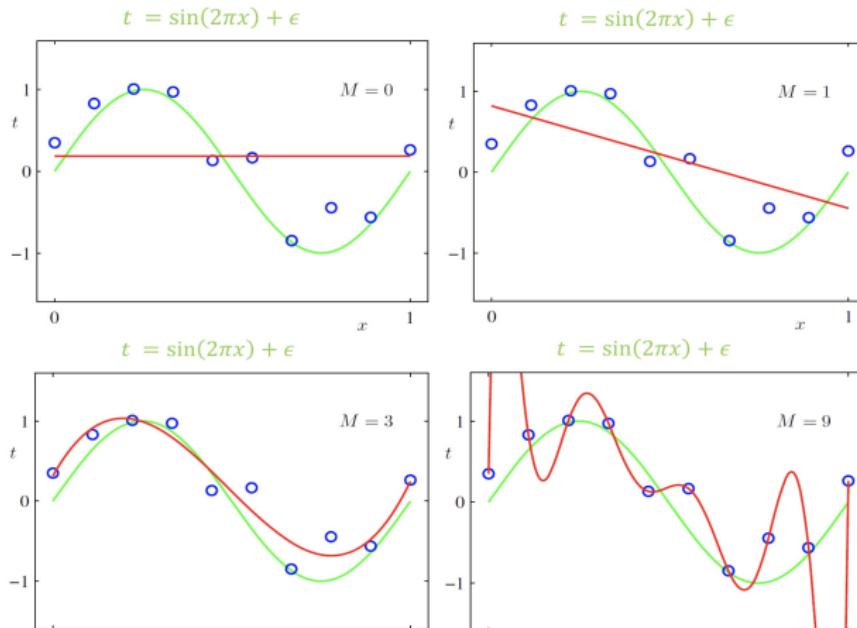
$$\begin{aligned} \frac{\partial}{\partial \mathbf{w}} (\mathbf{X}\mathbf{w} - \mathbf{y})^\top (\mathbf{X}\mathbf{w} - \mathbf{y}) + \lambda \tilde{\mathbf{w}}^\top \tilde{\mathbf{w}} &= \mathbf{0} \\ \Rightarrow 2\mathbf{X}^\top \mathbf{X}\mathbf{w} - 2\mathbf{X}^\top \mathbf{y} + 2\lambda \hat{\mathbf{I}}_d \mathbf{w} &= \mathbf{0} \\ \Rightarrow \mathbf{X}^\top \mathbf{X}\mathbf{w} + \lambda \hat{\mathbf{I}}_d \mathbf{w} &= \mathbf{X}^\top \mathbf{y} \\ \Rightarrow (\mathbf{X}^\top \mathbf{X} + \lambda \hat{\mathbf{I}}_d) \mathbf{w} &= \mathbf{X}^\top \mathbf{y} \\ \Rightarrow \mathbf{w} &= (\mathbf{X}^\top \mathbf{X} + \lambda \hat{\mathbf{I}}_d)^{-1} \mathbf{X}^\top \mathbf{y} \end{aligned}$$

Note that  $(\mathbf{X}^\top \mathbf{X} + \lambda \hat{\mathbf{I}}_d)$  is guaranteed to be invertible, given  $\lambda > 0$ .

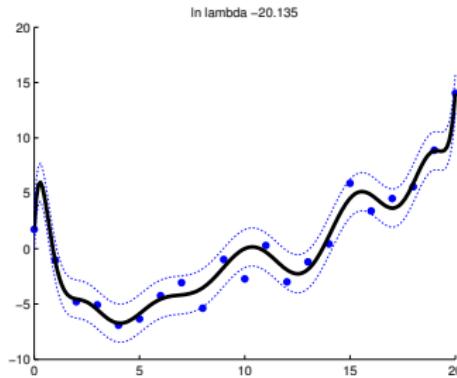
# Ridge regression

## Motivation 2:

- Overfitting is an important challenge for linear regression, as shown below.  
Note:  $M$  in the figure denotes the degree of polynomial hypothesis function.
- If overfitting, the prediction performance on testing data will be very poor.  
How to alleviate overfitting?



# Ridge regression



- Let's see one simple example, we use a **polynomial function** (introduced later) with 14 degrees to fit  $m = 21$  data points. The learned curve is very "wiggly" (see above).
- The parameter values of this curve are as follows
$$6.56, -36.934, -109.25, 543.452, 1022.561, -3046.224, -3768.013, 8524.54, 6607.897, -12640.058, -5530.188, 9479.73, 1774, 639, -2821.526$$
- There are many large positive/negative values, such that a small change of features could lead to a significant change of output.

# Ridge regression

- How to get smaller parameter values?
- We can assume that the parameter  $\mathbf{w}$  (excluding the bias  $b$ ) follow a zero-mean Gaussian prior

$$p(\mathbf{w}) = \mathcal{N}(\mathbf{w} | \mathbf{0}, \tau^2 \mathbf{I}) \quad (9)$$

- For clarity, we omit the bias  $w_0$  in the following derivation.
- Utilizing this prior, we obtain the maximum a posteriori (MAP) estimation

$$\mathbf{w}_{MAP} = \arg \max_{\mathbf{w}} \left[ \sum_{i=1}^m \log p(y_i | \mathbf{x}_i, \mathbf{w}) + \log p(\mathbf{w}) \right] \quad (10)$$

$$= \arg \max_{\mathbf{w}} \left[ \sum_{i=1}^m \log \mathcal{N}(\mathbf{x}_i^\top \mathbf{w}, \sigma^2) + \log \mathcal{N}(\mathbf{w} | \mathbf{0}, \tau^2 \mathbf{I}) \right] \quad (11)$$

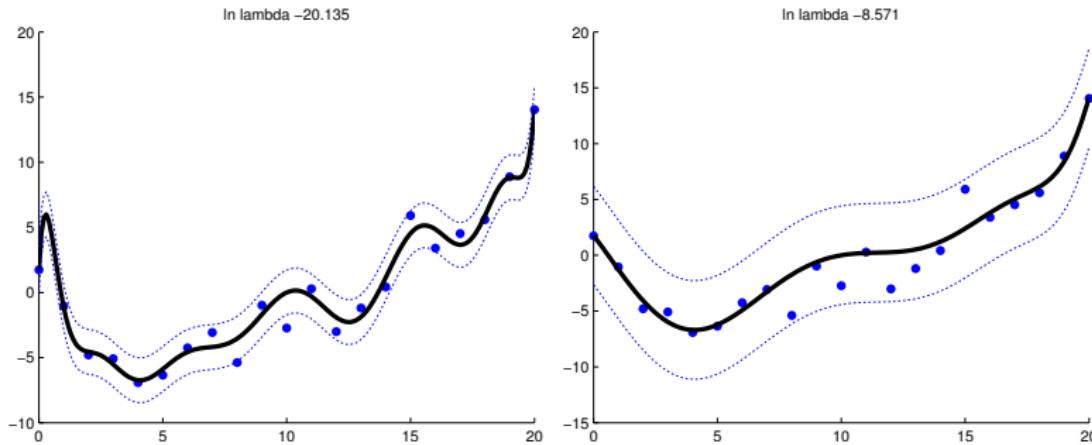
$$\equiv \arg \min_{\mathbf{w}} \left[ \sum_{i=1}^m (\mathbf{x}_i^\top \mathbf{w} - y_i)^2 + \lambda \|\mathbf{w}\|_2^2 \right]. \quad (12)$$

- The corresponding closed-form solution is given by

$$\mathbf{w}_{MAP} = (\lambda \mathbf{I} + \mathbf{X}^\top \mathbf{X})^{-1} \mathbf{X}^\top \mathbf{y}. \quad (13)$$

# Ridge regression

- The above method is also known as ridge regression, or **penalized least squares**.
- In general, adding a Gaussian prior to the parameters of a model to encourage them to be small is called  **$\ell_2$  regularization** or **weight decay**.
- As shown below, when we set a larger  $\lambda$ , i.e., more weight on the prior, the resulting curve will be smoother.

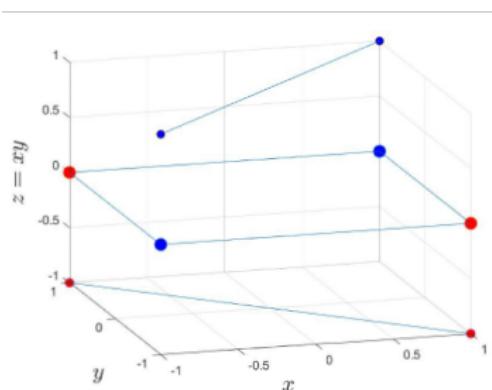
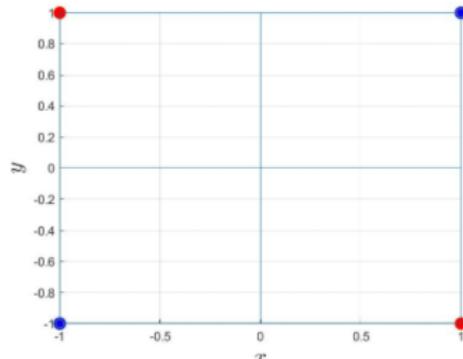


# Polynomial regression

## Motivation

- Some data may be not linearly separated, such as the classic **XOR** data, as shown on the bottom left.
- Consequently, the linear regression model doesn't work.
- To tackle it, we could project the original data to the **monomial** axis  $x_1x_2$ .
- Then, the XOR becomes linearly separated, as shown on the bottom right.
- Accordingly, we can design a novel linear regression model, as follows

$$f_{\mathbf{w}}(\mathbf{x}) = w_0 + w_1x_1 + w_2x_2 + \textcolor{green}{w_{12}x_1x_2} + w_{11}x_1^2 + w_{22}x_2^2$$



# Polynomial regression

## Polynomial expansion

- The linear model  $f_{\mathbf{w}}(\mathbf{x}) = \mathbf{x}^\top \mathbf{w}$  can be written as

$$f_{\mathbf{w}}(\mathbf{x}) = \sum_{i=0}^d x_i w_i = \textcolor{blue}{w_0} + \sum_{i=1}^d x_i w_i.$$

- By including terms involving the products of pairs of components of  $\mathbf{x}$ , we obtain a **quadratic** model:

$$f_{\mathbf{w}}(\mathbf{x}) = \textcolor{blue}{w_0} + \sum_{i=1}^d w_i x_i + \sum_{i=1}^d \sum_{j=1}^d w_{ij} x_i x_j.$$

# Polynomial regression

- In general:

$$f_{\mathbf{w}}(\mathbf{x}) = w_0 + \sum_{i=1}^d w_i x_i + \sum_{i=1}^d \sum_{j=1}^d w_{ij} x_i x_j + \sum_{i=1}^d \sum_{j=1}^d \sum_{k=1}^d w_{ijk} x_i x_j x_k + \dots$$

## Remarks

- For high dimensional  $d$  and high *polynomial order*, the number of polynomial terms becomes explosive! (In fact, this number of terms grows exponentially.)
- Hence, for high dimensional problems, polynomials of order larger than 3 are seldom used.

# Polynomial regression

Linear model with **basis expansion**  $\phi(\mathbf{x})$

$$\begin{aligned} f_{\mathbf{w}}(\mathbf{x}) &= \textcolor{blue}{w_0} + \sum_{i=1}^d w_i x_i + \sum_{i=1}^d \sum_{j=1}^d w_{ij} x_i x_j + \sum_{i=1}^d \sum_{j=1}^d \sum_{k=1}^d w_{ijk} x_i x_j x_k + \dots \\ &= \phi(\mathbf{x})^\top \mathbf{w}, \end{aligned}$$

where

$$\begin{aligned} \phi(\mathbf{x}) &= [1, x_1, \dots, x_d, \dots, x_i x_j, \dots, x_i x_j x_k, \dots]^\top, \\ \mathbf{w} &= [\textcolor{blue}{w_0}, w_1, \dots, w_d, \dots, w_{ij}, \dots, w_{ijk}, \dots]^\top. \end{aligned}$$

**Note:**  $f_{\mathbf{w}}(\mathbf{x})$  is still a **linear** function *w.r.t.*  $\mathbf{w}$ , rather than  $\mathbf{x}$ . Thus, it is still a **linear model**.

Extending to the case of  $m$  data points, *i.e.*,  $\mathbf{X} = [\mathbf{x}_1^\top; \dots; \mathbf{x}_m^\top] \in \mathbb{R}^{m \times (d+1)}$ , the basis expansion is presented by

$$\mathbf{P}(\mathbf{X}) = [\phi(\mathbf{x}_1)^\top; \dots; \phi(\mathbf{x}_m)^\top] \in \mathbb{R}^{m \times |\mathbf{w}|}.$$

# Polynomial regression

## Example

2<sup>nd</sup> order polynomial model

$$\begin{aligned}\{\mathbf{x}_i, \mathbf{y}_i\}_{i=1}^m & \quad \{x_1 = 0, x_2 = 0\} \rightarrow \{y = 0\} \\ & \quad \{x_1 = 1, x_2 = 1\} \rightarrow \{y = 1\} \\ & \quad \{x_1 = 1, x_2 = 0\} \rightarrow \{y = 2\} \\ & \quad \{x_1 = 0, x_2 = 1\} \rightarrow \{y = 3\}\end{aligned}$$

$$f_{\mathbf{w}}(\mathbf{x}) = w_0 + w_1 x_1 + w_2 x_2 + w_{12} x_1 x_2 + w_{11} x_1^2 + w_{22} x_2^2$$

$$= [1 \quad x_1 \quad x_2 \quad x_1 x_2 \quad x_1^2 \quad x_2^2] \begin{bmatrix} w_0 \\ w_1 \\ w_2 \\ w_{12} \\ w_{11} \\ w_{22} \end{bmatrix}$$

$$\mathbf{P} = \begin{bmatrix} 1 & x_1 & x_2 & x_1 x_2 & x_1^2 & x_2^2 \\ 1 & x_1 & x_2 & x_1 x_2 & x_1^2 & x_2^2 \\ 1 & x_1 & x_2 & x_1 x_2 & x_1^2 & x_2^2 \\ 1 & x_1 & x_2 & x_1 x_2 & x_1^2 & x_2^2 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 & 1 \end{bmatrix}$$

# Polynomial regression

Ridge regression with **original features  $\mathbf{X}$** :

$$\begin{aligned}\text{Learning: } \hat{\mathbf{w}} &= (\mathbf{X}^\top \mathbf{X} + \lambda \mathbf{I})^{-1} \mathbf{X}^\top \mathbf{y} \\ \text{Prediction: } f_{\mathbf{w}}(\mathbf{X}_{new}) &= \mathbf{X}_{new} \hat{\mathbf{w}}\end{aligned}$$

Ridge regression with **basis expansion  $\mathbf{P}(\mathbf{X})$** :

$$\begin{aligned}\text{Learning: } \hat{\mathbf{w}} &= (\mathbf{P}^\top \mathbf{P} + \lambda \mathbf{I})^{-1} \mathbf{P}^\top \mathbf{y} \\ \text{Prediction: } f_{\mathbf{w}}(\mathbf{P}(\mathbf{X}_{new})) &= \mathbf{P}_{new} \hat{\mathbf{w}}\end{aligned}$$

# Polynomial regression

## For Regression Applications

- Learning:  $\hat{\mathbf{w}} = (\mathbf{P}^\top \mathbf{P} + \lambda \mathbf{I})^{-1} \mathbf{P}^\top \mathbf{y}$ , where  $\mathbf{y}$  is continuous
- Prediction:  $f_{\mathbf{w}}(\mathbf{P}(\mathbf{X}_{new})) = \mathbf{P}_{new} \hat{\mathbf{w}}$

## For Classification Applications

- Learn **discrete** valued  $\mathbf{y}$  (binary) or  $\mathbf{Y}$  (one-hot)
- Binary Prediction:  $f_{\mathbf{w}}(\mathbf{P}(\mathbf{X}_{new})) = \text{sgn}(\mathbf{P}_{new} \hat{\mathbf{w}})$  if  $y \in \{-1, +1\}$
- Multi-Category Prediction:  $f_{\mathbf{w}}(\mathbf{P}(\mathbf{X}_{new})) = \text{argmax}_{i=1, \dots, C} (\mathbf{P}_{new} \hat{\mathbf{w}})$

# Polynomial regression

## Example (Cont'd)

$$\begin{aligned}\{x_i, y_i\}_{i=1}^m \quad & \{x_1 = 0, x_2 = 0\} \rightarrow \{y = -1\} \\ & \{x_1 = 1, x_2 = 1\} \rightarrow \{y = -1\} \\ & \{x_1 = 1, x_2 = 0\} \rightarrow \{y = +1\} \\ & \{x_1 = 0, x_2 = 1\} \rightarrow \{y = +1\}\end{aligned}$$

2<sup>nd</sup> order polynomial model

$$P = \begin{bmatrix} 1 & x_1 & x_2 & x_1 x_2 & x_1^2 & x_2^2 \\ 1 & x_1 & x_2 & x_1 x_2 & x_1^2 & x_2^2 \\ 1 & x_1 & x_2 & x_1 x_2 & x_1^2 & x_2^2 \\ 1 & x_1 & x_2 & x_1 x_2 & x_1^2 & x_2^2 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 & 0 & 0 & 0 \\ 1 & 1 & 1 & 1 & 1 & 1 \\ 1 & 1 & 0 & 0 & 1 & 0 \\ 1 & 0 & 1 & 0 & 0 & 1 \end{bmatrix}$$

$$\hat{w} = P^T (P P^T)^{-1} y \quad (\text{note: under determined linear system})$$

$$= \begin{bmatrix} 1 & 1 & 1 & 1 \\ 0 & 1 & 1 & 0 \\ 0 & 1 & 0 & 1 \\ 0 & 1 & 0 & 0 \\ 0 & 1 & 1 & 0 \\ 0 & 1 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 1 & 1 & 1 \\ 1 & 6 & 3 & 3 \\ 1 & 3 & 3 & 1 \\ 1 & 3 & 1 & 3 \end{bmatrix}^{-1} \begin{bmatrix} -1 \\ -1 \\ +1 \\ +1 \end{bmatrix} = \begin{bmatrix} -1 \\ 1 \\ 1 \\ -4 \\ 1 \\ 1 \end{bmatrix}$$

# Polynomial regression

## Example (Cont'd)

Prediction:

Test point 1:  $\{x_1 = 0.1, x_2 = 0.1\} \rightarrow \{y = \text{class } -1 \text{ or } +1?\}$

Test point 2:  $\{x_1 = 0.9, x_2 = 0.9\} \rightarrow \{y = \text{class } -1 \text{ or } +1?\}$

Test point 3:  $\{x_1 = 0.1, x_2 = 0.9\} \rightarrow \{y = \text{class } -1 \text{ or } +1?\}$

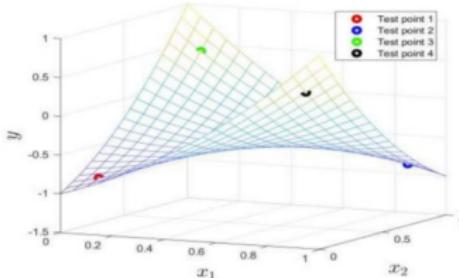
Test point 4:  $\{x_1 = 0.9, x_2 = 0.1\} \rightarrow \{y = \text{class } -1 \text{ or } +1?\}$

$$\hat{y} = \mathbf{P}_t \hat{\mathbf{w}}$$

$$\hat{y} = \text{sgn}\left(\begin{bmatrix} 1 & 0.1 & 0.1 & 0.01 & 0.01 & 0.01 \\ 1 & 0.9 & 0.9 & 0.81 & 0.81 & 0.81 \\ 1 & 0.1 & 0.9 & 0.09 & 0.01 & 0.81 \\ 1 & 0.9 & 0.1 & 0.09 & 0.81 & 0.01 \end{bmatrix} \begin{bmatrix} -1 \\ 1 \\ 1 \\ -4 \\ 1 \\ 1 \end{bmatrix}\right)$$

$$= \text{sgn}\left(\begin{bmatrix} -0.82 \\ -0.82 \\ 0.46 \\ 0.46 \end{bmatrix}\right)$$

$$= \begin{bmatrix} -1 \\ -1 \\ 1 \\ 1 \end{bmatrix} \begin{array}{l} \xrightarrow{\quad} \text{Class } -1 \\ \xrightarrow{\quad} \text{Class } -1 \\ \xrightarrow{\quad} \text{Class } +1 \\ \xrightarrow{\quad} \text{Class } +1 \end{array}$$



# Lasso regression

- We can replace the Gaussian prior by a Laplacian prior, i.e.,

$$p(\mathbf{w}) = \text{Lap}(\mathbf{w}|\mathbf{0}, \lambda) = \frac{1}{2\lambda} \exp\left(-\frac{\|\mathbf{w}\|_1}{\lambda}\right), \quad (14)$$

- The combination of the Gaussian distribution of  $p(y|\mathbf{x}, \mathbf{w})$  and the Laplacian prior, leading to

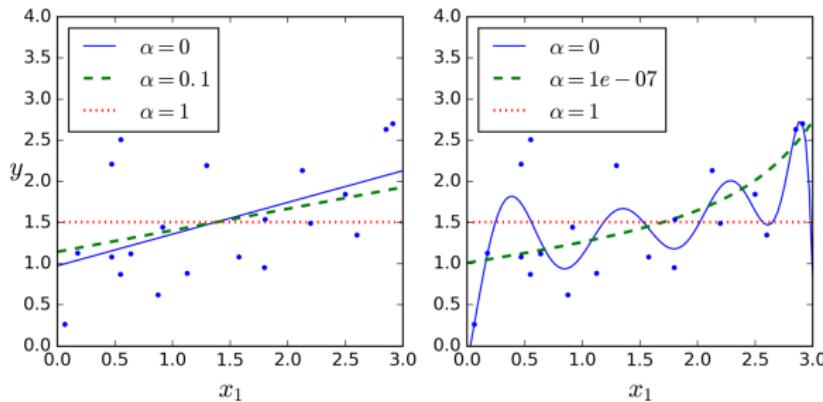
$$\mathbf{w}_{MAP} = \arg \max_{\mathbf{w}} \left[ \sum_{i=1}^m \log p(y_i|\mathbf{x}_i, \mathbf{w}) + \log p(\mathbf{w}) \right] \quad (15)$$

$$= \arg \max_{\mathbf{w}} \left[ \sum_{i=1}^m \log \mathcal{N}(\mathbf{w}^\top \mathbf{x}_i, \sigma^2) + \text{Lap}(\mathbf{w}|\mathbf{0}, b) \right] \quad (16)$$

$$\equiv \arg \min_{\mathbf{w}} \left[ \sum_{i=1}^m (\mathbf{x}_i^\top \mathbf{w} - y_i)^2 + \alpha \|\mathbf{w}\|_1 \right]. \quad (17)$$

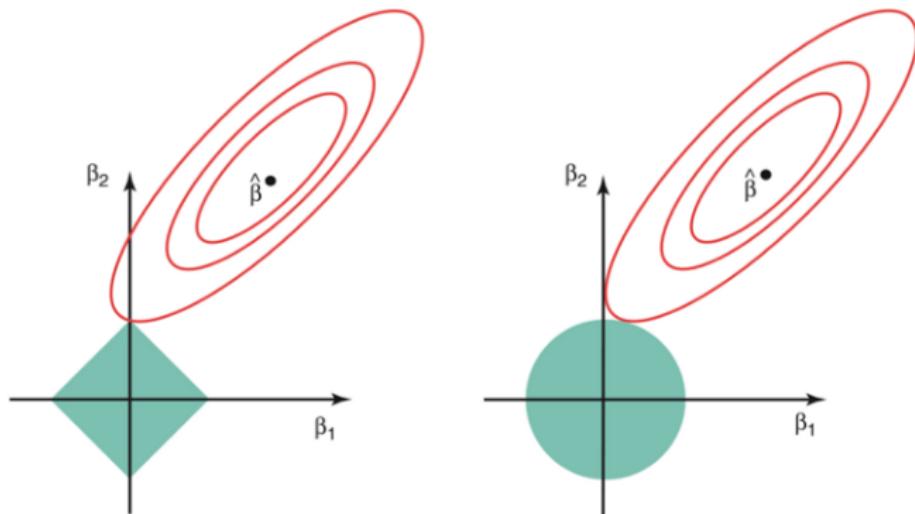
# Lasso regression

- It is **Lasso regression**, and the regularization is called  $\ell_1$  regularization.  
It will encourage the sparse parameters.
- As shown below, when we set a larger  $\alpha$ , *i.e.*, more weight on the prior, the resulting curve will be smoother.



# Geometry of Ridge and Lasso regression

- Geometry of Ridge and Lasso regression. Which one is Ridge?

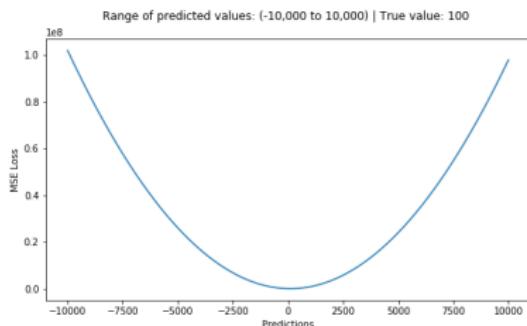
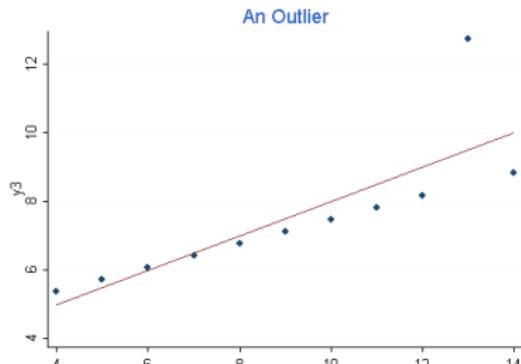


# Robust linear regression

- When there is a few outliers in the training data  $D$ , which are far from most other points, then learned parameters  $\mathbf{w}_{MLE}$  will be significantly influenced, leading to a very poor fitting.
- Let's see the loss curve of the residual sum of squares (RSS),

$$J(\mathbf{w}) = \frac{1}{2} \sum_{i=1}^m (\mathbf{x}_i^\top \mathbf{w} - y_i)^2. \quad (18)$$

- The error increases quadratically along with the residual. To minimize such a large error, the linear model will be significantly changed.
- How to alleviate the significant influence of outliers?

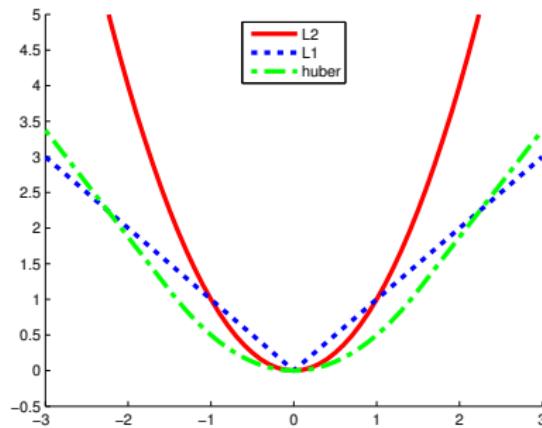


# Robust linear regression

- We adopt the  $\ell_1$  loss to replace the  $\ell_2$  loss, as follows

$$J(\mathbf{w}) = \sum_{i=1}^m |\mathbf{x}_i^\top \mathbf{w} - y_i|. \quad (19)$$

- The curves of  $\ell_1$  and  $\ell_2$  losses are shown as follows.
- When the residual is large, the  $\ell_1$  loss is **much smaller** than the  $\ell_2$  loss, such that the influence of outliers could be alleviated.



# Robust linear regression

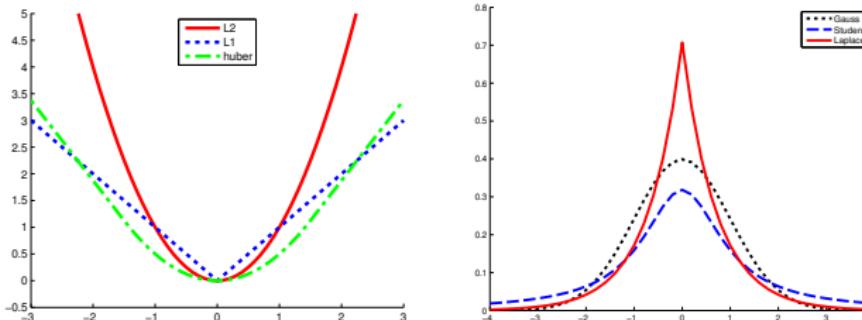
- Actually, the above  $\ell_1$  loss can also be derived from the probabilistic perspective, by assuming that

$$p(y|\mathbf{x}, \mathbf{w}, b) = \text{Lap}(y|\mathbf{w}^\top \mathbf{x}, b) \propto \exp\left(-\frac{1}{b}|y - \mathbf{w}^\top \mathbf{x}|\right) \quad (20)$$

- Applying the maximum log-likelihood estimation (MLE), we will obtain

$$\mathbf{w}_{MLE} = \arg \max_{\mathbf{w}} \log \mathcal{L}(\mathbf{w}; D) = \arg \max_{\mathbf{w}} \sum_i^m \log p(y_i|\mathbf{x}_i, \mathbf{w}) \quad (21)$$

$$\equiv \arg \min_{\mathbf{w}} \frac{1}{b} \sum_{i=1}^m |\mathbf{x}_i^\top \mathbf{w} - y_i| \quad (22)$$



# Robust linear regression

$$\mathbf{w}_{MLE} = \arg \min_{\mathbf{w}} \sum_{i=1}^m |\mathbf{x}_i^\top \mathbf{w} - y_i| \quad (23)$$

- However, the  $\ell_1$  loss function is **non-differentiable**. The gradient descent algorithm cannot be adopted.
- We can transform it to a **linear program**, as follows

$$\begin{aligned} & \min_{\mathbf{w}, \mathbf{t}} \sum_i^m \mathbf{t}_i \\ & \text{s.t. } -\mathbf{t}_i \leq \mathbf{x}_i^\top \mathbf{w} - y_i \leq \mathbf{t}_i, 1 \leq i \leq m. \end{aligned} \quad (24)$$

Please refer to:

<https://math.stackexchange.com/questions/1639716/how-can-l-1-norm-minimization-with>

# Robust linear regression

$$\mathbf{w}_{MLE} = \arg \min_{\mathbf{w}} \sum_{i=1}^m |\mathbf{x}_i^\top \mathbf{w} - y_i| \quad (25)$$

- We can also utilize the following equation:

$$|a| = \min_{\mu > 0} \frac{1}{2} \left( \frac{a^2}{\mu} + \mu \right) \quad (26)$$

- Then, the above  $\ell_1$  minimization problem (25) can be reformulated as follows

$$\min_{\mathbf{w}} \min_{\mu_1, \dots, \mu_m > 0} \frac{1}{2} \sum_{i=1}^m \left( \frac{(\mathbf{x}_i^\top \mathbf{w} - y_i)^2}{\mu_i} + \mu_i \right). \quad (27)$$

- It can be iteratively and alternatively optimized as follows:
  - Given  $\mathbf{w}$ ,  $\mu_i = |\mathbf{x}_i^\top \mathbf{w} - y_i|$ ,  $i = 1, \dots, m$
  - Given  $\mu$ ,  $\mathbf{w} = \arg \min_{\mathbf{w}} \sum_{i=1}^m \frac{1}{2} (\mathbf{x}_i^\top \mathbf{w} - y_i)^2 / \mu_i$
- It is called **iteratively reweighted least squares** method.

# Summary of different variants of linear regressions

Note that the uniform distribution will not change the mode of the likelihood.

Thus, MAP estimation with a uniform prior corresponds to MLE.

| $p(y \mathbf{x}, \mathbf{w})$ | $p(\mathbf{w})$ | regression method |
|-------------------------------|-----------------|-------------------|
| Gaussian                      | Uniform         | Least squares     |
| Gaussian                      | Gaussian        | Ridge regression  |
| Gaussian                      | Laplace         | Lasso regression  |
| Laplace                       | Uniform         | Robust regression |
| Student                       | Uniform         | Robust regression |

